

Aha: Moments

as meta-cognitive prediction errors

Hermish Mehta University of California, Berkeley



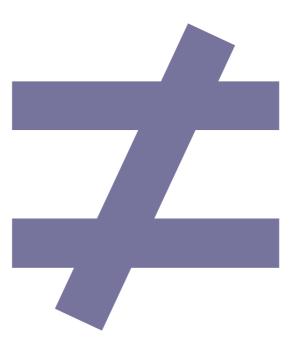
Ana. Moments

as meta-cognitive prediction errors

in collaboration with Rachit Dubey, Mark Ho, Tom Griffiths Princeton University

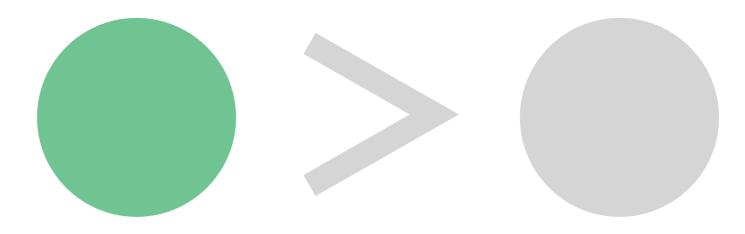
outline

- 1. Motivation
- 2. Background
- 3. Theoretical Model
- 4. Experimental Results
- 5. Parting Thoughts



Not all inferences are equally rewarding.

Motivation



Why are some inferences more rewarding than others?

questions

- 1. Why do we have *Aha!* moments?
- 2. What makes these moments so rare?
- 3. Why are they so rewarding?

computational algorithmic implementational

What is the goal of the system? How can we formulate the objective?

algorithmic

implementational

What is the goal of the system? How can we formulate the objective?

What is an appropriate algorithm to achieve this goal?

implementational

What is the goal of the system? How can we formulate the objective?

What is an appropriate algorithm to achieve this goal?

How is the algorithm implemented?



computational algorithmic implementational How? (physical)

Why? (problem)

What? (rules)



computational

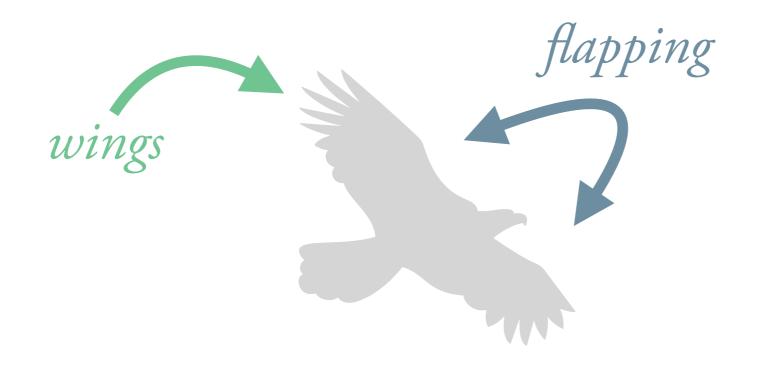
algorithmic

implementational

Why? (problem)

What? (rules)

How? (physical)



computational

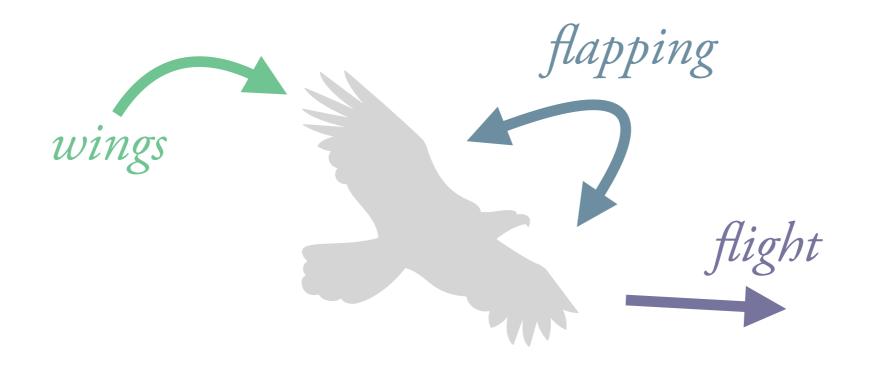
algorithmic

implementational

Why? (problem)

What? (rules)

How? (physical)



computational
algorithmic
implementational

Why? (problem)

What? (rules)

How? (physical)

"Trying to understand perception by studying only neurons is like trying to understand bird flight by studying only feathers: It just cannot be done."



key idea

Aha! moments arise under resource constraints and are meta-cognitive prediction errors which cause reward prediction errors.

meta-cognitive prediction errors

reward prediction errors

resource constraints

Background

learning l'larning

(noun)

A non-temporary change in the behavioral mechanisms engaged in a certain situation... that results from repeated experience with the situation... and providing the change can't be explained in terms of innate behavioral tendencies.

learning l'larning

The process of improving predictions of the future.

```
    {prediction error}
    =
    {actual outcome outcome outcome}
    -
    {predicted outcome outcome}
```

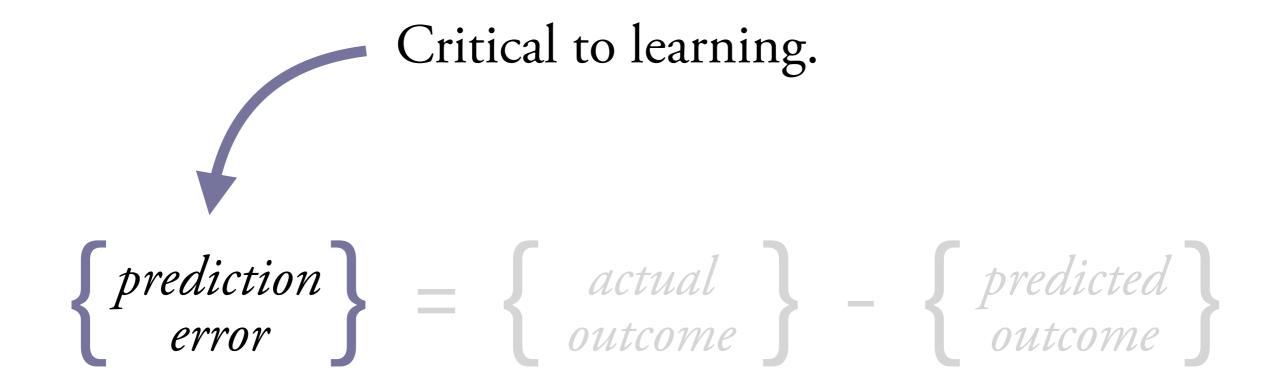
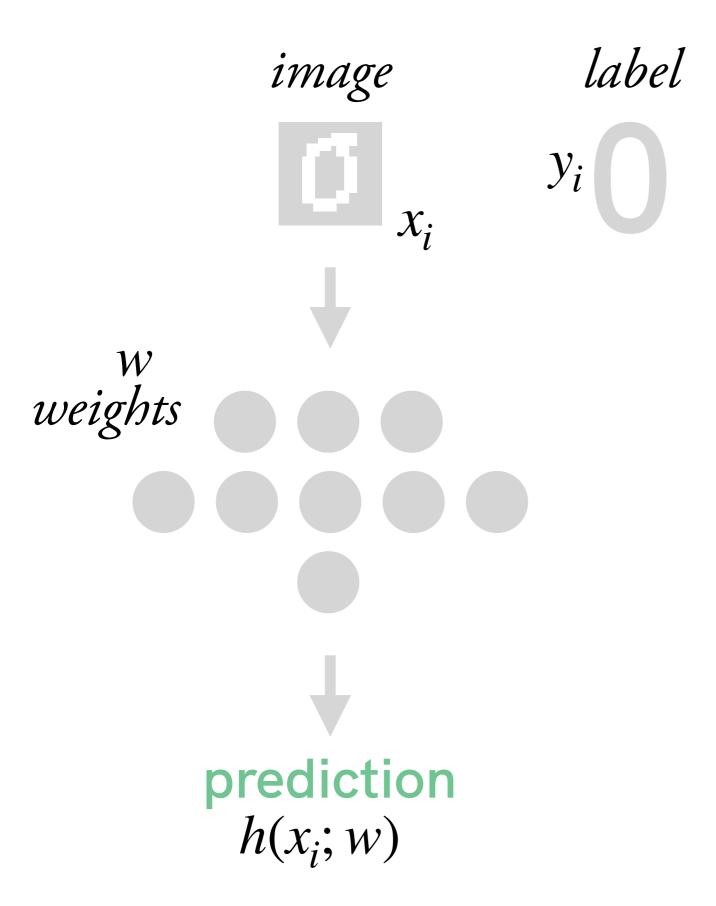
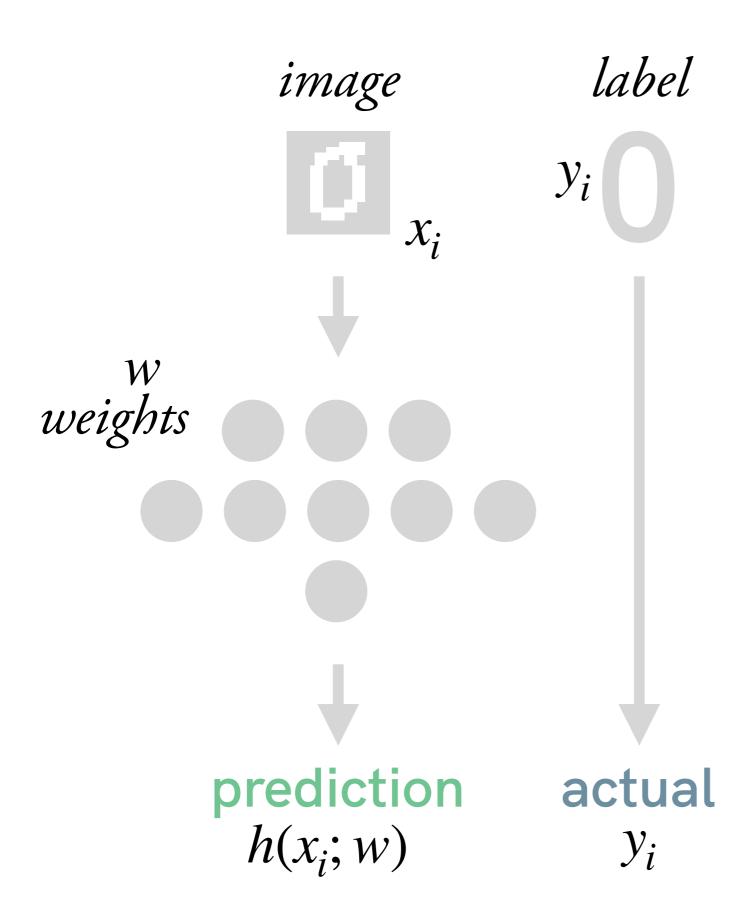
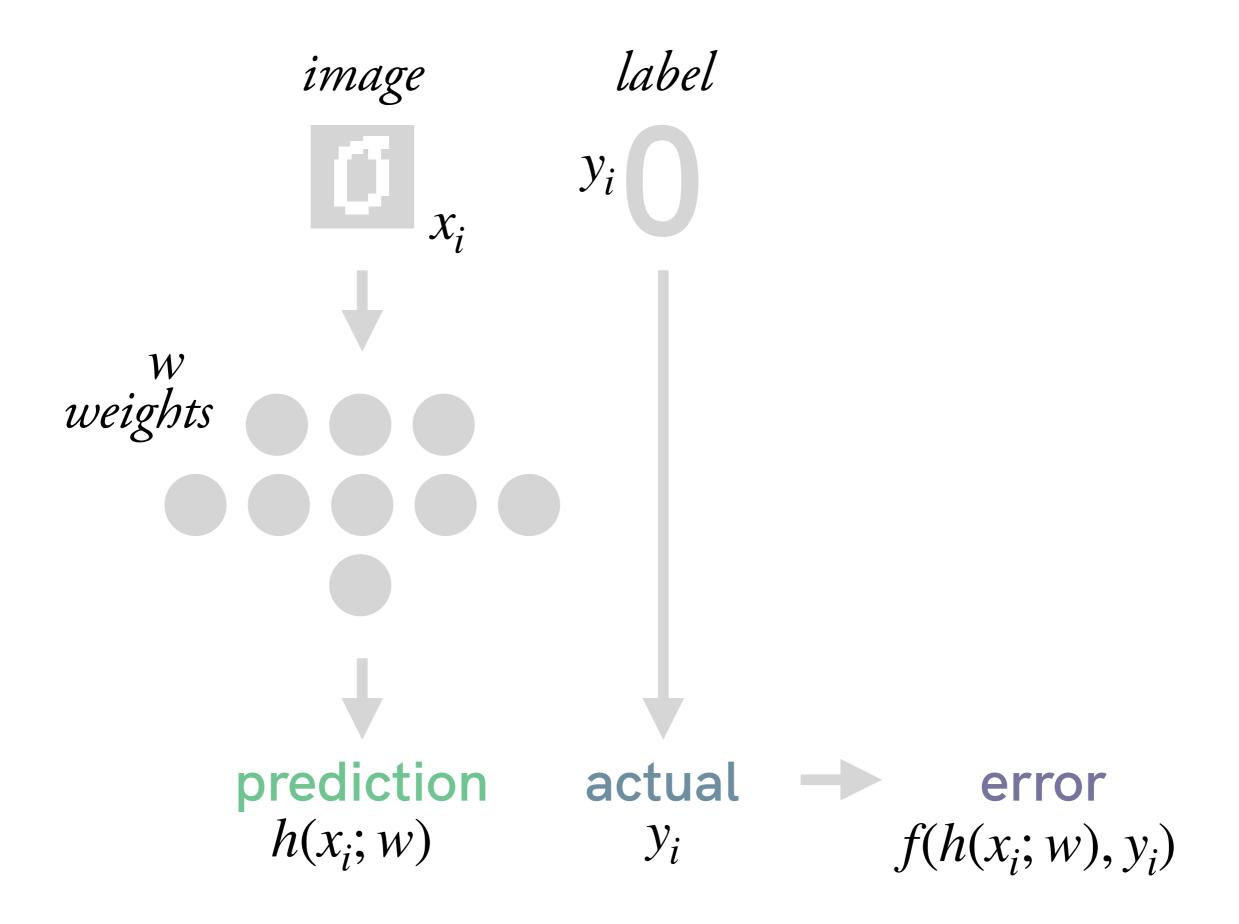


image label

 X_i Y_i







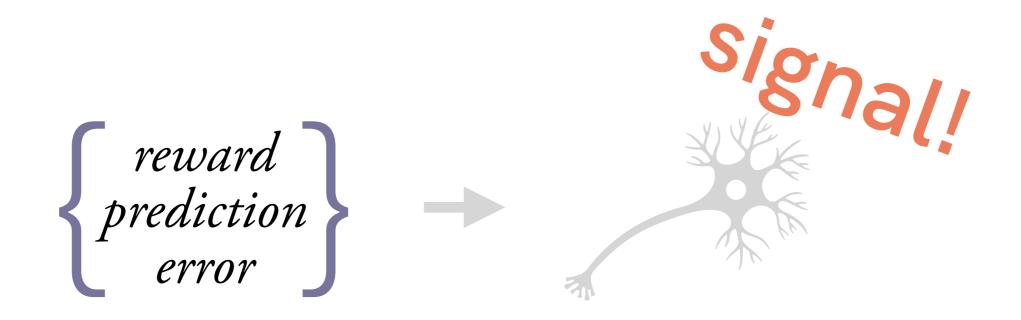
machine learning

$$w^* = \arg\min_{w} \frac{1}{n} \sum_{i=1}^{n} f(h(x_i; w), y_i)$$

Minimize average prediction error.

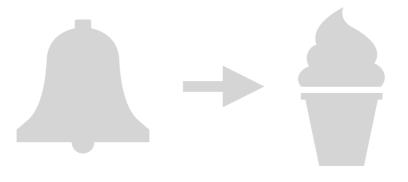


Rewards can also be predicted.

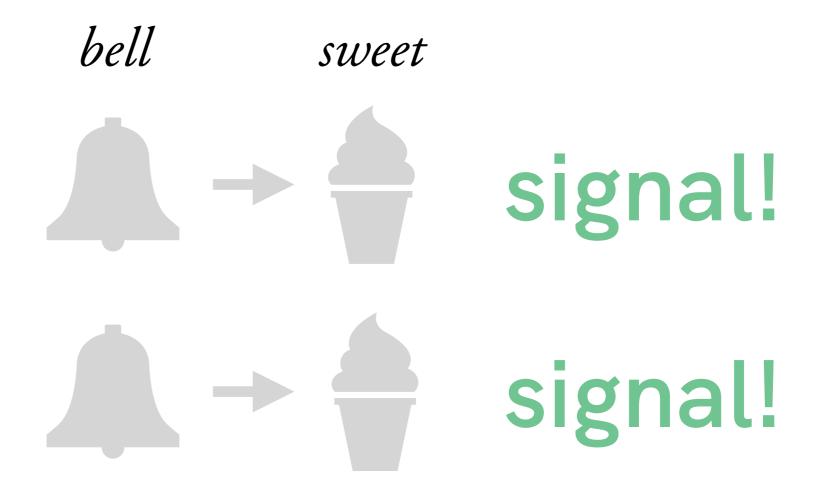


The phasic dopamine response signals prediction error.

bell sweet



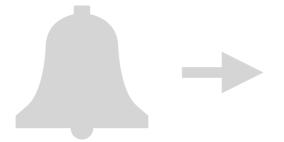
bell sweet Signal!



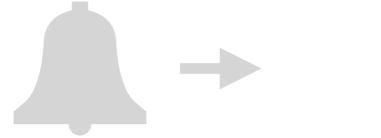
bell sweet signal! signal!



bell

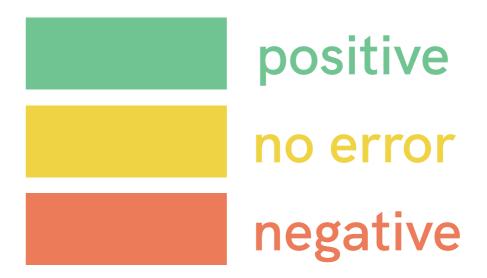


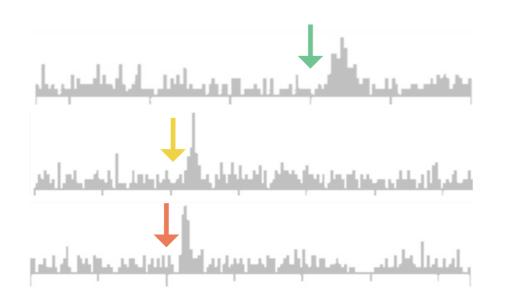
bell





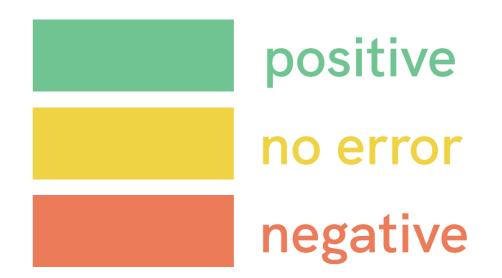
reward prediction error





dopamine response

reward prediction error



Bayer, Hannah M., and Paul W. Glimcher. "Midbrain dopamine neurons encode a quantitative reward prediction error signal." *Neuron* 47.1 (2005): 129-141.

Hammer, Martin. "The neural basis of associative reward learning in honeybees." *Trends in neurosciences* 20.6 (1997): 245-252.

Hollerman, Jeffrey R., and Wolfram Schultz. "Dopamine neurons report an error in the temporal prediction of reward during learning." *Nature neuroscience* 1.4 (1998): 304.

Schultz, Wolfram, Peter Dayan, and P. Read Montague. "A neural substrate of prediction and reward." *Science* 275.5306 (1997): 1593-1599.

Hart, Andrew S., et al. "Phasic dopamine release in the rat nucleus accumbens symmetrically encodes a reward prediction error term." *Journal of Neuroscience* 34.3 (2014): 698-704.

Colombo, Matteo. "Deep and beautiful. The reward prediction error hypothesis of dopamine." Studies in history and philosophy of science part C: Studies in history and philosophy of biological and biomedical sciences 45 (2014): 57-67.

bounded optimality

Artificial agents need to tradeoff between efficiency and accuracy.

computational limitations & costs



algorithm 1 algorithm 2 algorithm 3 algorithm 4

resource rationality

People need to make optimal use of their limited resources, which involves "meta-reasoning," thinking about thinking.

Griffiths, Thomas L., Falk Lieder, and Noah D. Goodman. "Rational use of cognitive resources: Levels of analysis between the computational and the algorithmic." *Topics in cognitive science* 7.2 (2015): 217-229.

Shenhav, Amitai, et al. "Toward a rational and mechanistic account of mental effort." *Annual review of neuroscience* 40 (2017): 99-124.

Lieder, Falk, Ming Hsu, and Thomas L. Griffiths. "The high availability of extreme events serves resource-rational decision-making." *Proceedings of the annual meeting of the cognitive science society.* Vol. 36. No. 36. 2014.

Lieder, Falk, and Thomas L. Griffiths. "Strategy selection as rational metareasoning." *Psychological Review* 124.6 (2017): 762.

Lieder, Falk, et al. "The anchoring bias reflects rational use of cognitive resources." *Psychonomic bulletin & review* 25.1 (2018): 322-349.

Theoretical Model

tasks

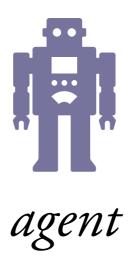












tasks

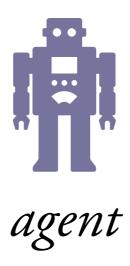






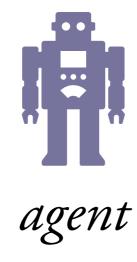






goal

Maximize reward.



tasks

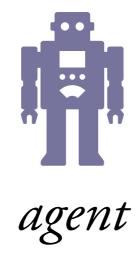






goal

Complete as many tasks as possible.



tasks

1 R

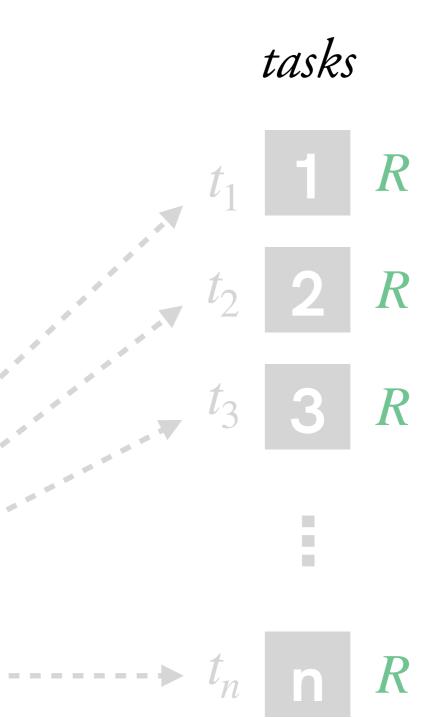
2 R

3 *R*

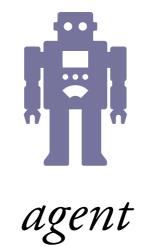
n R

goal

Complete as many tasks as possible.

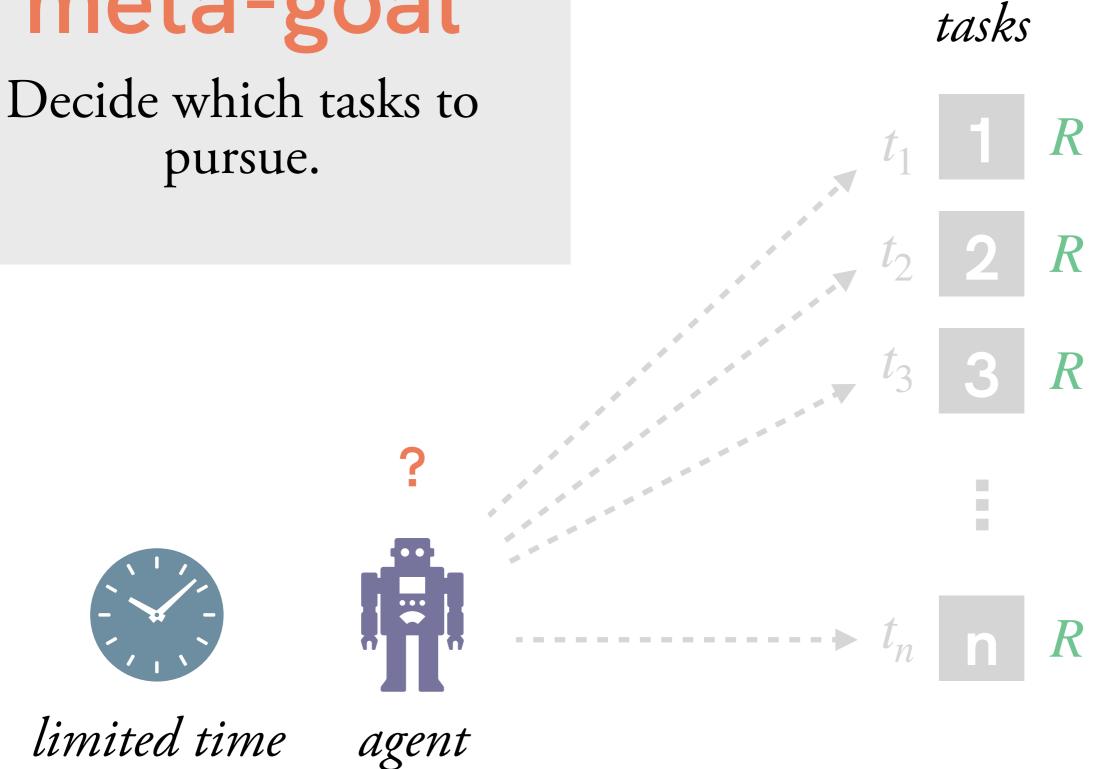






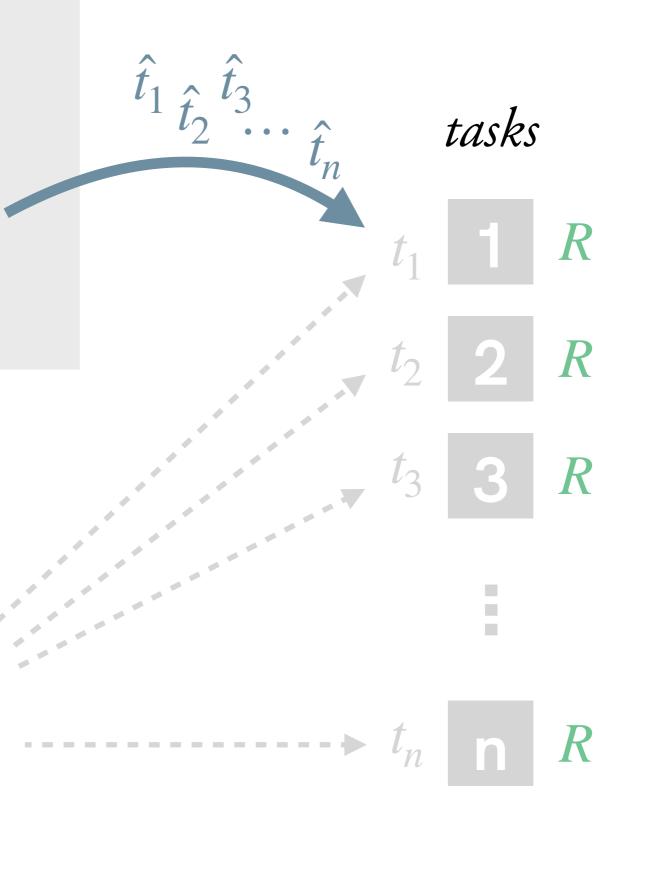
meta-goal

pursue.



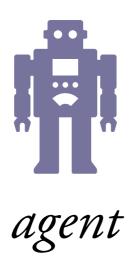
solution

Estimate completion times.







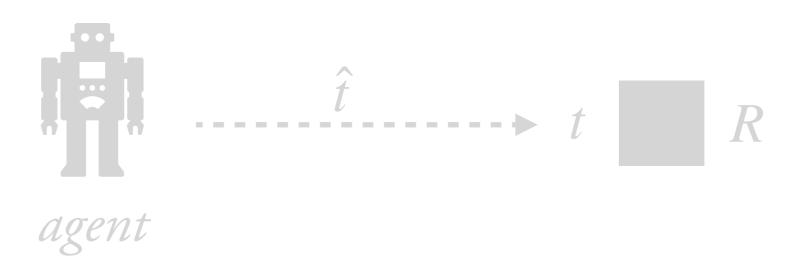


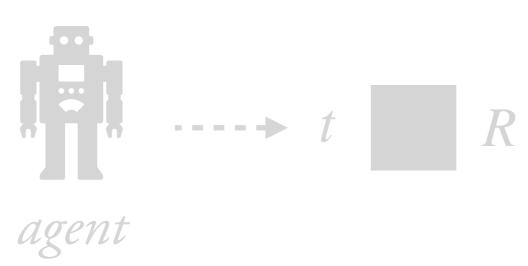




(meta-cognitive) time prediction error

$$\hat{t} > t$$
 positive $\hat{t} = t$ no error $\hat{t} < t$ negative





 $\hat{t} > t$

positive reward prediction error





 $\hat{t} > t$

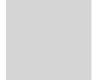
aha!



positive reward prediction error

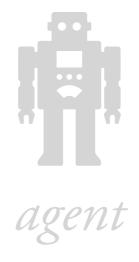






$\hat{t} < t$

negative rewa prediction er

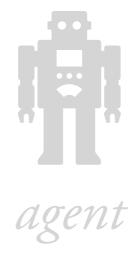






$\hat{t} < t$

gah!



negative rewa prediction err



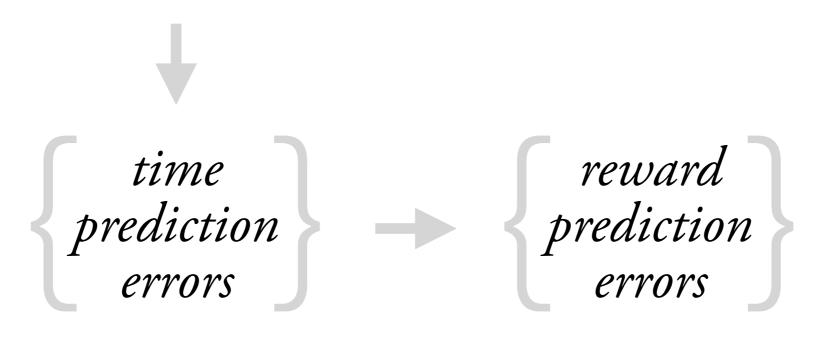
constraints time predictions

time predictions

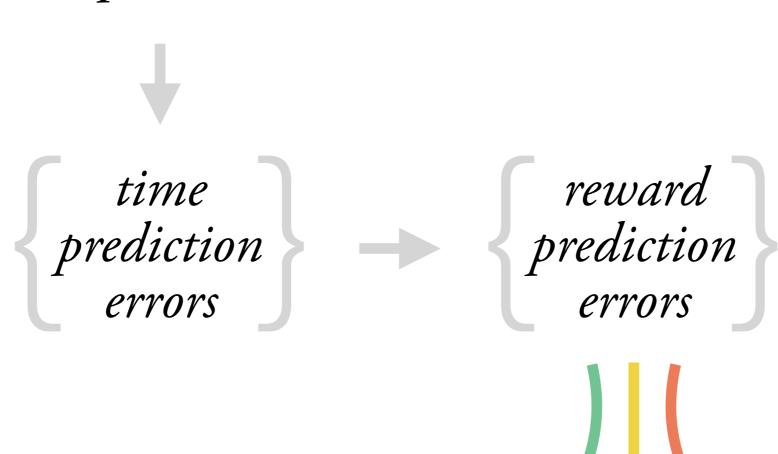


time prediction errors

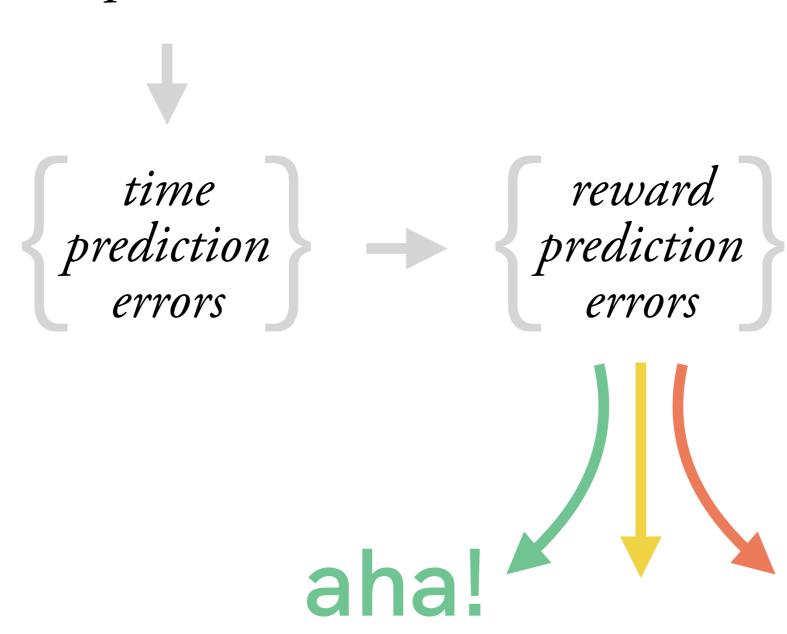
time predictions



time predictions

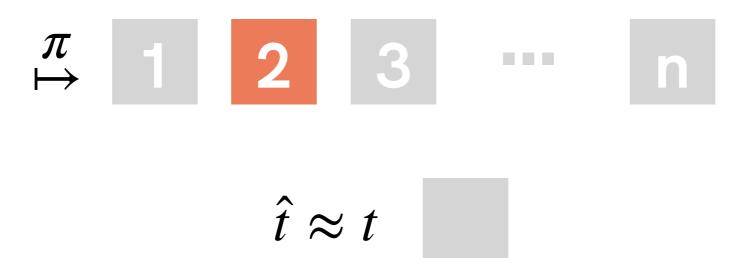


time predictions

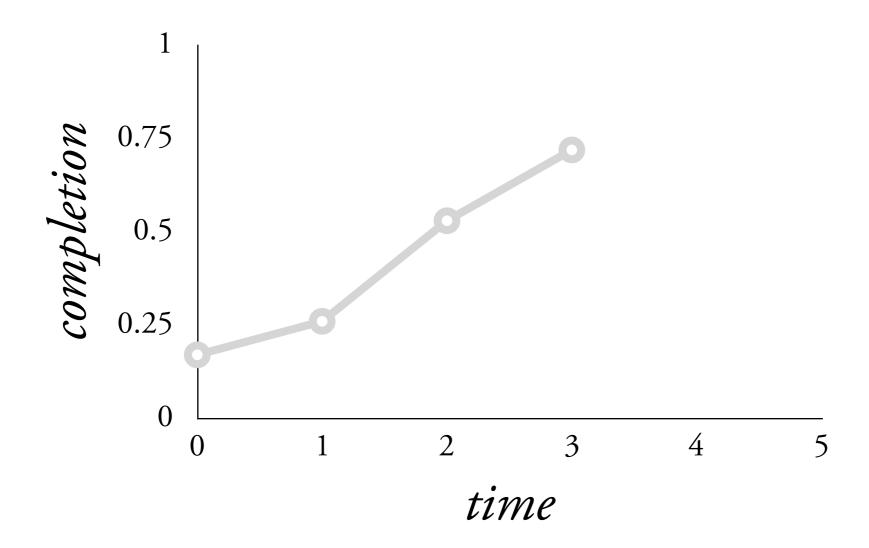


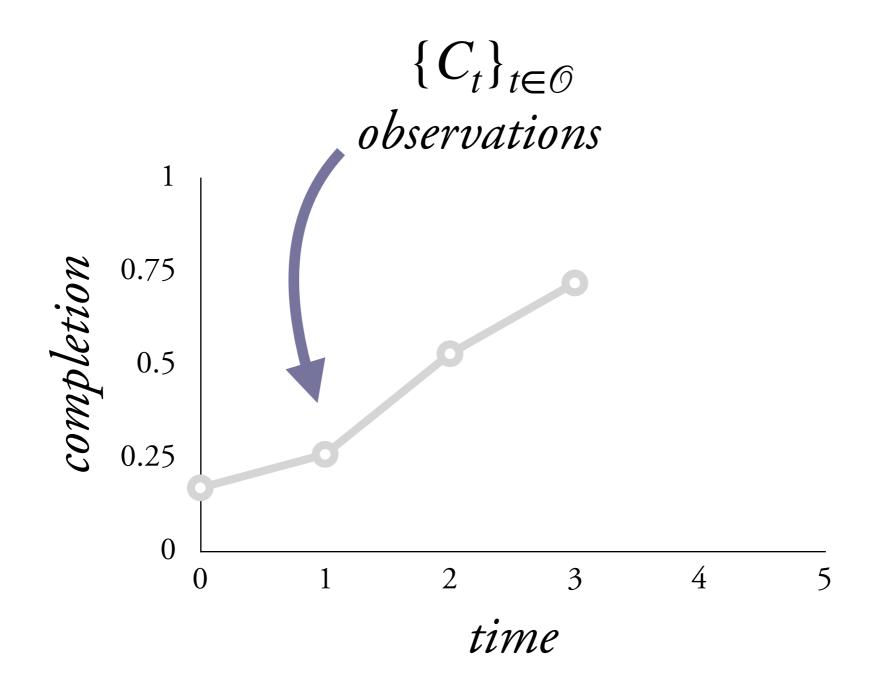
key idea

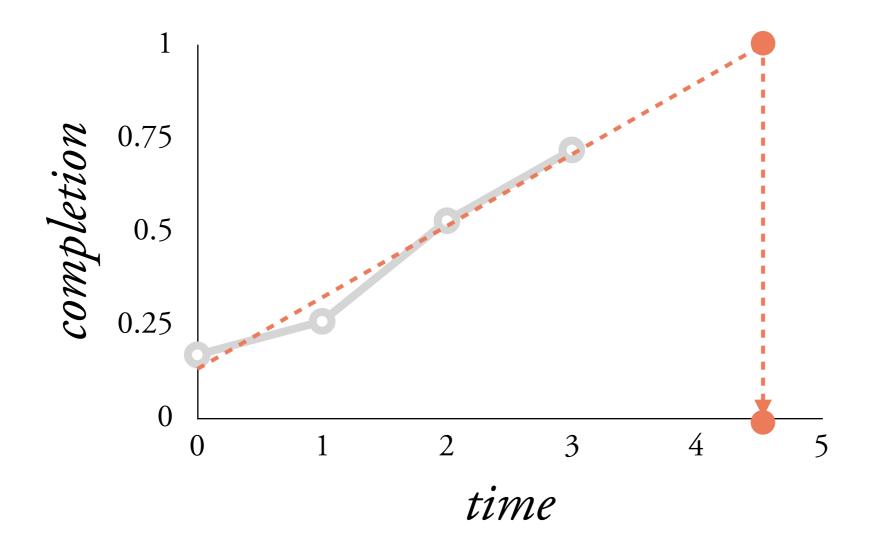
Aha! moments arise under resource constraints and are meta-cognitive prediction errors which cause reward prediction errors.



Our framework leaves out two essential pieces: task selection and time estimation.







step completion

$$d_1$$
 $c_1 = 0.05$
 d_2 $c_2 = 0.15$
 d_3 $c_3 = 0.24$

 $d_k \quad c_k = 0.67$

$$d_1$$
 $c_1 = 0.05$

$$d_2$$
 $c_2 = 0.15$

$$d_3$$
 $c_3 = 0.24$

•

$$d_k \quad c_k = 0.67$$



$$X \sim f_X(\cdot;\theta)$$
.

$$d_1$$
 $c_1 = 0.05$

$$d_2$$
 $c_2 = 0.15$

$$d_3$$
 $c_3 = 0.24$

•

$$d_k \quad c_k = 0.67$$

Understanding increases by some random variable

$$X \sim f_X(\cdot;\theta)$$
.

goal

Estimate time to completion.

bayesian estimate

probability given parameters and history...



$$\Pr(C_m \ge 1) = \int_{\mathbb{R}^d} \Pr(C_m \ge 1 \mid \theta, \mathbf{C} = \mathbf{c}) \cdot f_{\mathbf{\Theta} \mid \mathbf{C}}(\theta \mid \mathbf{c}) \, \mathrm{d}\theta$$



probability completed at step m



...weighted by posteriors

predictions

- 1. Aha! moments emerge when an impasse is suddenly broken.
- 2. Aha! moments do not emerge when learning occurs gradually.

$$d_1$$
 $c_1 = 0.01$
 d_2 $c_2 = 0.05$
 d_3 $c_3 = 0.05$
 \vdots

 $d_k \quad c_k = 0.06$

$$d_1$$
 $c_1 = 0.01$

$$d_2$$
 $c_2 = 0.05$

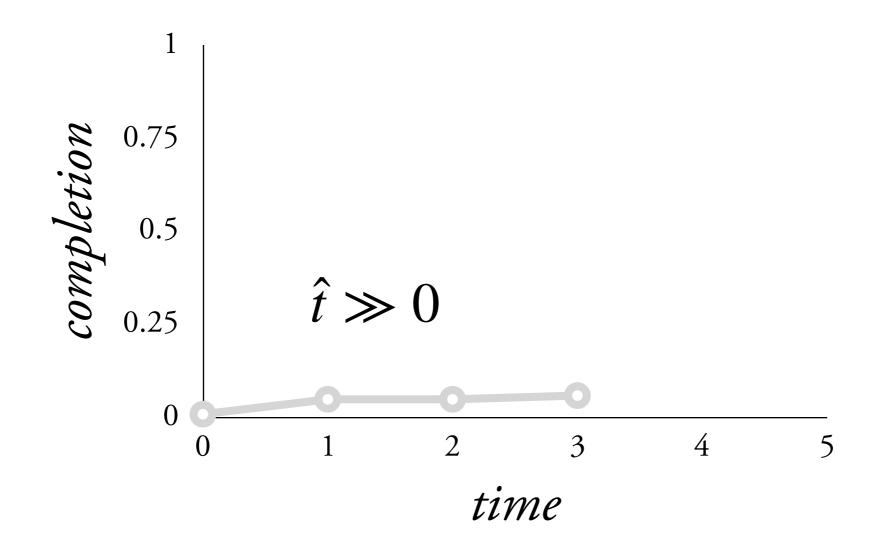
$$d_3$$
 $c_3 = 0.05$

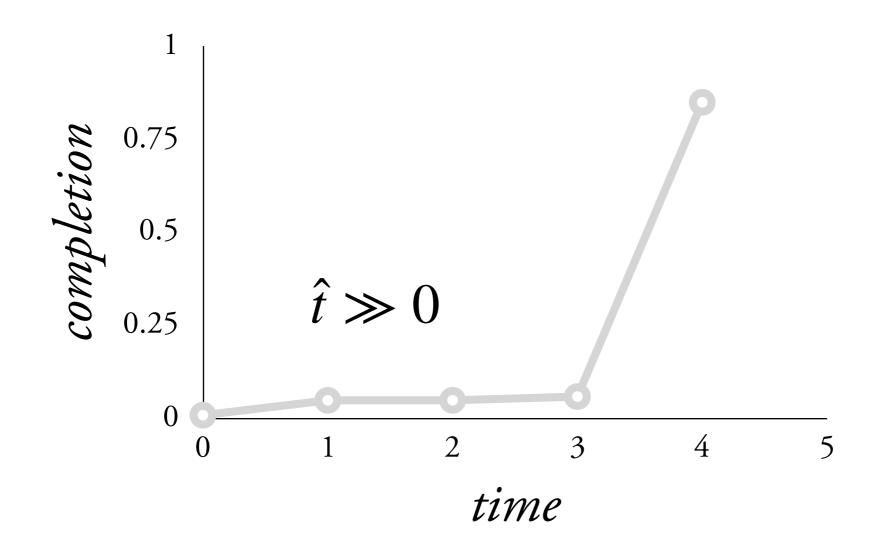
•

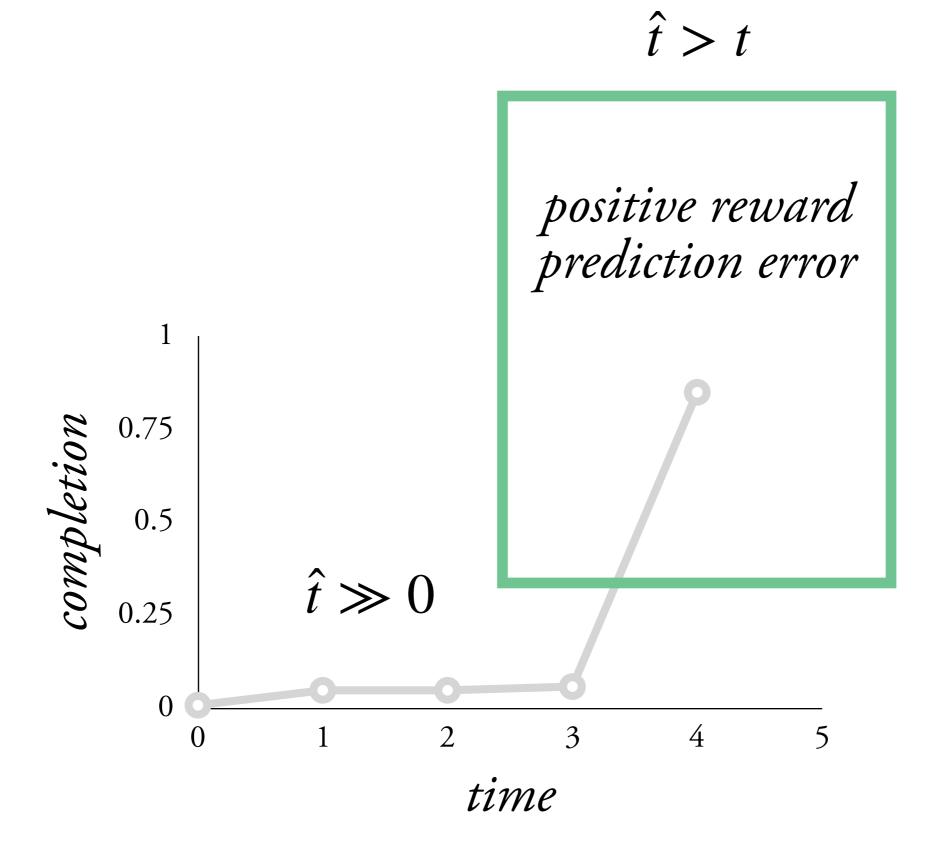
$$d_k c_k = 0.06$$

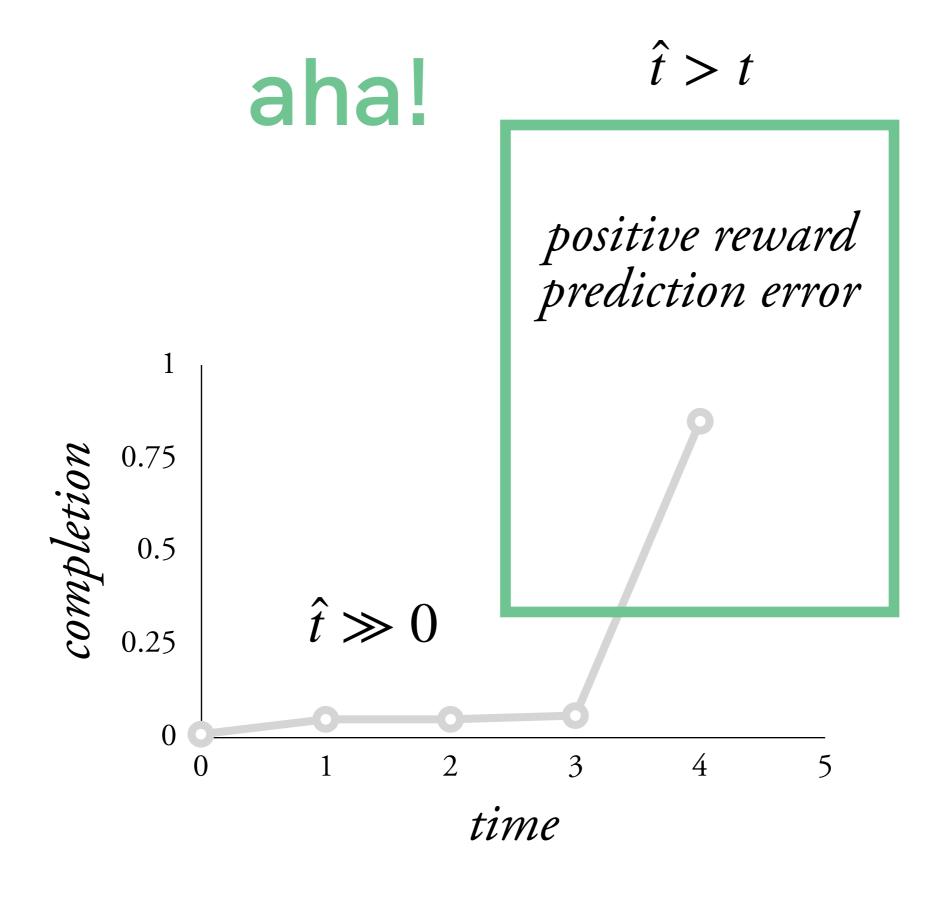
estimate

$$\hat{t} \gg 0$$









$$d_1$$
 $c_1 = 0.05$
 d_2 $c_2 = 0.15$
 d_3 $c_3 = 0.30$

 $d_k c_k = 0.4$

$$d_1$$
 $c_1 = 0.05$

$$d_2$$
 $c_2 = 0.15$

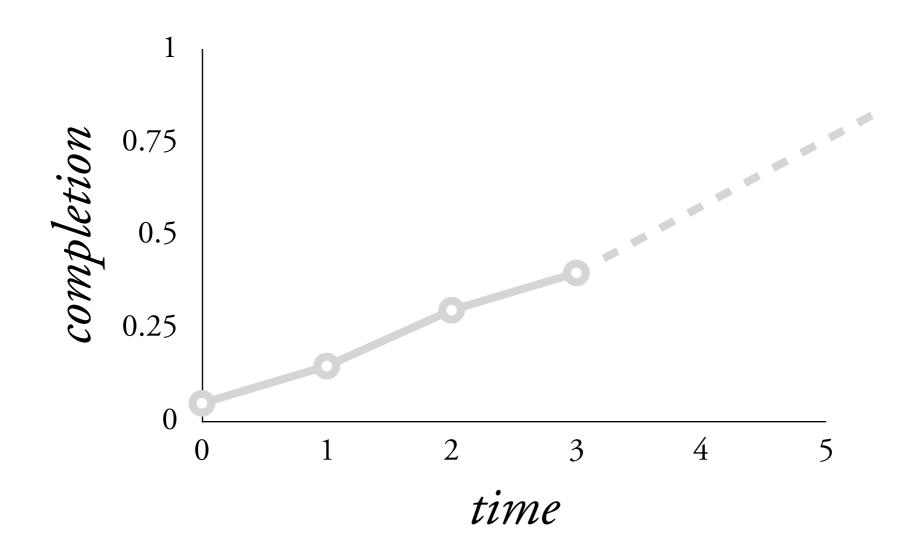
$$d_3$$
 $c_3 = 0.30$

•

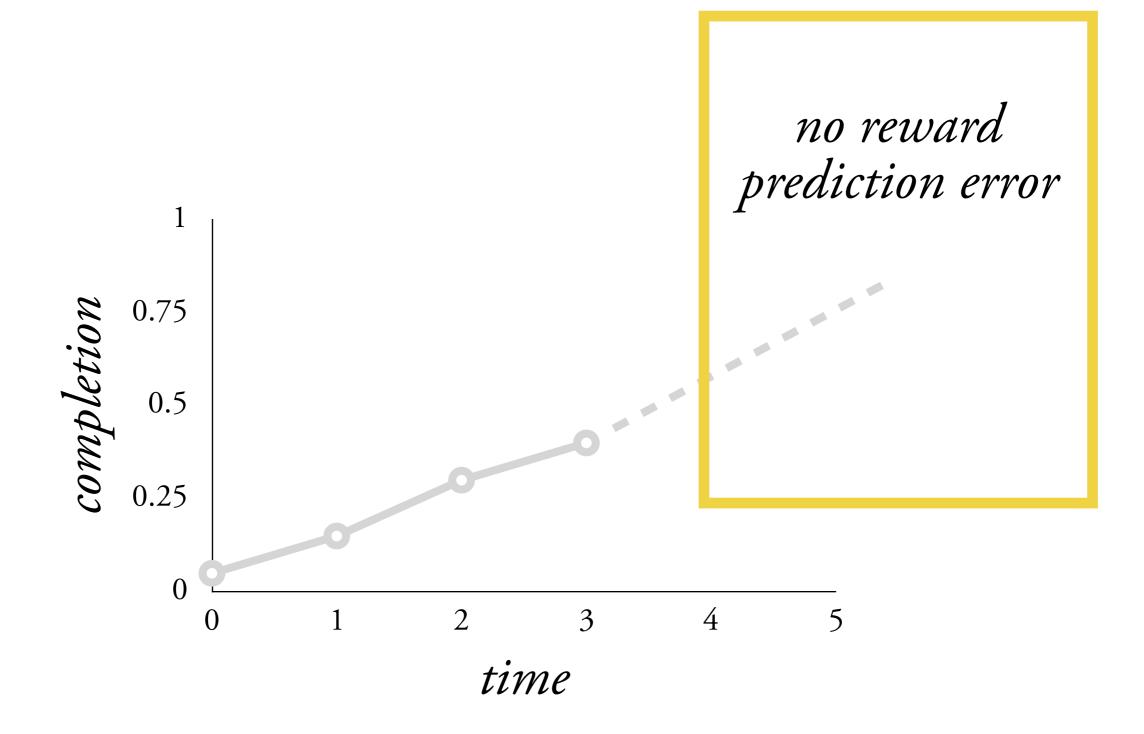
$$d_k c_k = 0.4$$

estimate

$$\hat{t} \approx t$$







other results

3. *Aha!* moments will be influenced by our prior expectations and are a product of our subjective experience.

other results

4. *Aha!* moments will be highest after we have spent considerable time on a problem and have been stuck on it for a while.

Experimental Results

question

Can solely manipulating people's perceived difficulty of a task influences their experience of an *Aha!* moment?

Participants were given 10 seconds for each of 5 anagrams.

SROT

anagram 1 of 5

Each anagram was immediately followed with an *Aha!* self-report.

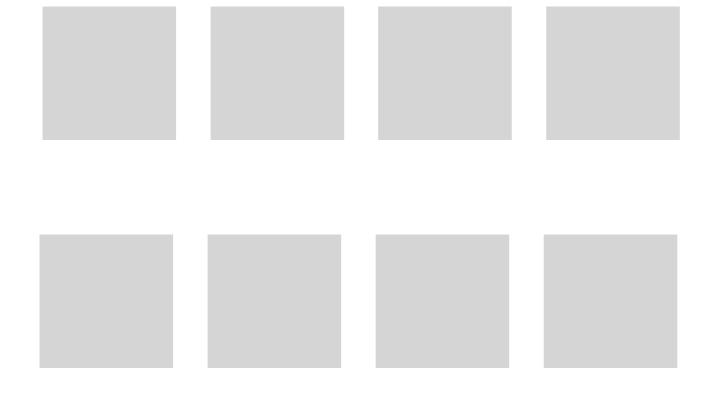
Did you experience an *Aha!* moment (on a scale of 1-7) after the **last** anagram was shown to you?

anagram 1 of 5

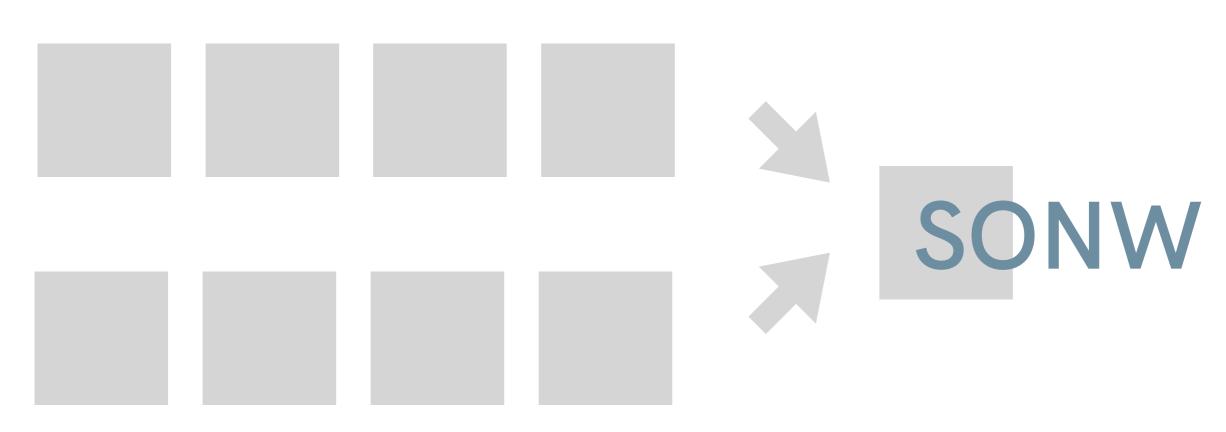
124
participants

64
easy group

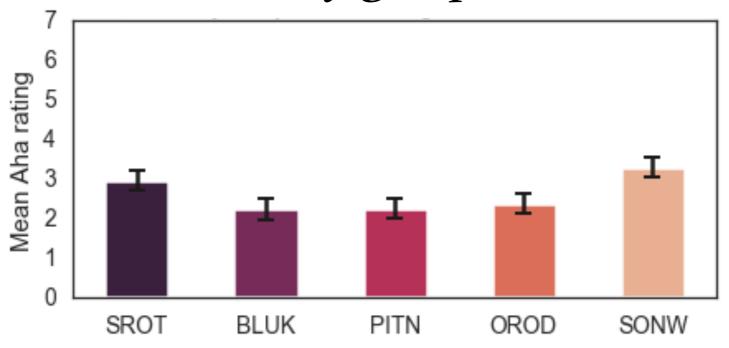
64
easy group

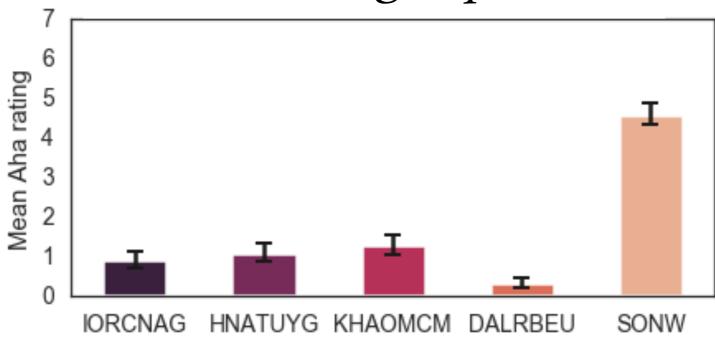


64
easy group

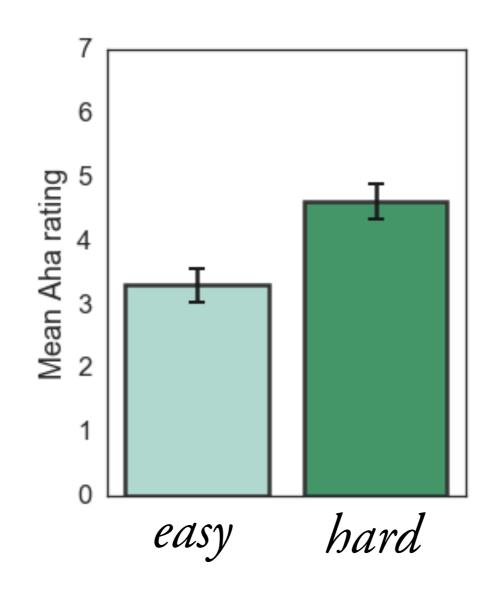


easy group





$$(t = 3.3, p < 0.01)$$



Parting Thoughts

questions

- 1. Why do we have *Aha!* moments?
- 2. What makes these moments so rare?
- 3. Why are they so rewarding?



2. What makes these moments so rare? *Positive time prediction errors are rare* (planning fallacy).



further work

We are currently working on more robust behavioral experiments to test our theory.



Aha! Moments

as meta-cognitive prediction errors

Ana. Moments

as meta-cognitive prediction errors

in collaboration with Rachit Dubey, Mark Ho, Tom Griffiths Princeton University University of California, Berkeley